Shreyam Kanaujiya

Software Engineer • Java • Spring Boot • Microservices • LLMs • RAG • React
Pune, MH | **shreyam91183@gmail.com** | +91 **6392764845** | **LinkedIn** | **GitHub** | **Portfolio**

Professional Summary

Software Engineer with **4+ years of experience** developing full-stack applications using **Java, Spring Boot, React, REST APIs, and microservices**. Experienced in backend performance optimization, database query tuning, frontend performance, and cloud deployments. Building AI applications with **Python, LLMs, RAG, and vector search**, including document retrieval and engineering-focused AI tools.

Technical Skills

Languages: Java, Python, SQL, JavaScript

AI & Generative AI: RAG, AI Agents, LLMs, Embeddings, Vector Search, Reranking, Prompt Engineering

Frontend: React, Next.js, Redux Toolkit

Backend: Spring Boot, FastAPI, REST APIs, Microservices

Databases: PostgreSQL, MySQL, Redis, pgvector

Cloud & DevOps: AWS (EC2, ECR), Docker, Git, GitHub Actions, CI/CD
Testing: JUnit, Unit Testing, API Testing, Postman

Professional Experience

Senior Software Engineer **Jun 2024 – Present**
Techium Infotech Pvt. Ltd. – Pune

- Migrated **8 microservices to cloud-based authentication,** refactoring legacy authentication and service-to-service authorization to support secure deployments without downtime.
- Analyzed Hibernate SQL queries to identify **N+1 fetch issues** and **database bottlenecks,** then optimized data fetching and introduced batch operations in Spring services, reducing average **API response time by 35%** under representative workloads.
- Configured **Resilience4j circuit breakers,** timeouts, and bounded retries to handle downstream service failures, **reducing dependency-related API errors by 30%** during transient disruptions.
- Reduced React application **LCP from 4.2 seconds to 2.1 seconds** through rendering optimization, route-level code splitting, lazy loading, and asset optimization.
- Built reusable React components and shared frontend utilities, reducing duplicated UI logic and development time by **30% across application modules.**
- Migrated **10 microservices** to **AWS** using **Docker, GitHub Actions, AWS ECR, and EC2,** automating build and deployment workflows and reducing deployment time by **60–70%.**
- Used AI tools to generate **unit tests, troubleshoot code, and automate repetitive development tasks,** reducing manual effort in routine coding and testing activities.

Junior Software Engineer **Sep 2022 – May 2024**
Cognizant Technology Solutions – Pune

- Built and integrated **15+ Spring Boot REST APIs across healthcare microservices,** implementing authentication, validation, business logic, exception handling, and service communication.
- Developed React.js interfaces for **healthcare registration and patient-data workflows, reducing dashboard load time from 6 seconds to 4 seconds** through API optimization, memoization, lazy loading, pagination, and component restructuring.
- Created Spring Boot services for a **Claims Management System,** implementing claim validation, processing, approval or rejection, and status tracking. Containerized services with **Docker** and deployed application components on **AWS.**
- Automated processing and segregation of **10,000+ Excel records using Java and Apache POI,** reducing manual effort by **90% and saving 30–40 minutes per file.**
- Improved troubleshooting by adding application logging and exception handling for API failures, validation errors, and service integration issues. Validated workflows using **100+ JUnit tests and Postman.**

Personal Projects

OpsMind — Secure Enterprise RAG & AI Assistant | Python, FastAPI, PostgreSQL, pgvector

- Built a document-based question-answering application using **PDF/text ingestion, chunking, embeddings, vector search, and reranking** to retrieve relevant information for LLM-generated responses.
- Implemented **database-level role-based access control (RBAC)** in vector search, along with source citations and prompt-injection safeguards to help prevent unauthorized or unsupported information from entering responses.
- Added **response evaluation and SSE streaming** to assess answer quality and deliver responses incrementally.

DevHub — Engineering Intelligence Platform | Next.js, GitHub, RAG, LLMs, pgvector

- Developed an agentic workflow that continuously tracks repository and documentation changes, analyzes new code and pull requests, and updates architecture documentation when engineering decisions or implementation details change.
- Developed LLM-based pull request analysis to flag potential bugs, architecture violations, missing test coverage, and conflicts with documented engineering decisions.
- Implemented RAG-based project intelligence using embeddings, pgvector, and metadata-aware retrieval, improving **relevant-context retrieval by 25%** across engineering queries.

Education

Bachelor of Technology in Computer Science **Aug 2018 – Jul 2022**
Pranveer Singh Institute of Technology

Certification

Agentic AI Foundation Associate Oracle